

Capstone 2 - Predicting Water Pump Condition EDA

September 8, 2019

1 Capstone 2 - Predicting Water Pump Condition in Tanzania EDA

Kenneth Liao

1.1 Background

The UN publishes and reviews a list of least developed countries (LDC) every 3 years. LDCs are “low-income countries confronting severe structural impediments to sustainable development. They are highly vulnerable to economic and environmental shocks and have low levels of human assets.”¹. Tanzania has been classified as an LDC since the UN published the first list of LDCs in 1971². A common challenge of LDCs is a lack of infrastructure to support the development of the nation, including access to education and healthcare, waste management, and access to potable water.

According to UNICEF, as of 2017, more than 24 million Tanzanians lacked access to basic drinking water³. This corresponds to only 56.7% of the country’s population having access to basic drinking water. Outside of developed urban areas, much of the potable water is accessed via water pumps.

Taarifa is an open-source platform for crowd-sourced reporting and triaging of infrastructure related issues. Together with the Tanzanian Ministry of Water, data has been collected for thousands of water pumps throughout Tanzania. The goal of this project is to be able to predict the condition of these water pumps to improve maintenance, reduce pump downtime, and ensure basic water access for millions of Tanzanians.

References

1. <https://www.un.org/development/desa/dpad/least-developed-country-category.html>
2. https://www.un.org/development/desa/dpad/wp-content/uploads/sites/45/publication/ldc_list.pdf
3. <https://washwatch.org/en/countries/tanzania/summary/statistics/>

1.1.1 Problem Description

Predict the operating condition of water pumps in Tanzania given various metadata on each water pump.

1.1.2 Strategy

The strategy will be to implement an XGBoost model as well as a neural network model for predictions and compare their performance.

1.1.3 Data

The dataset is provided by Taarifa, together with the Tanzanian Ministry of Water and is hosted by DrivenData.org:

<https://www.drivendata.org/competitions/7/pump-it-up-data-mining-the-water-table/page/23/>

1.2 Exploratory Data Analysis

Start by importing the necessary libraries and datasets.

```
In [1]: import pandas as pd
import plotly.graph_objs as go
from plotly.offline import iplot, plot, init_notebook_mode
import plotly.express as px
from config import credentials

init_notebook_mode(connected=True)

In [2]: # load the data
train = pd.read_csv('../data/train.csv')
train_labels = pd.read_csv('../data/train-labels.csv')
```

1.2.1 Prediction Labels

```
In [3]: train_labels.shape
```

```
Out[3]: (59400, 2)
```

```
In [4]: train_labels.head()
```

```
Out[4]:
```

	id	status_group
0	69572	functional
1	8776	functional
2	34310	functional
3	67743	non functional
4	19728	functional

```
In [5]: train_labels.id.nunique()
```

```
Out[5]: 59400
```

The train_labels file contains the labels we want to predict, status_group. This is the condition of a given water pump.

```
In [6]: counts = train_labels.groupby('status_group').count()
counts
```

```
Out [6]:
```

	id
status_group	
functional	32259
functional needs repair	4317
non functional	22824

```
In [7]: trace0 = go.Bar(name='functional', x=['functional'], y=[counts.loc['functional','id']]
        marker=dict(color='lightgreen'), showlegend=False)

        trace1 = go.Bar(name='functional needs repair', x=['functional needs repair'], y=[counts.loc['functional needs repair','id']]
        marker=dict(color='orange'), showlegend=False)

        trace2 = go.Bar(name='non functional', x=['non functional'], y=[counts.loc['non functional','id']]
        marker=dict(color='tomato'), showlegend=False)

        layout = go.Layout(title='Pump Condition Distribution',
        yaxis=dict(title='Count'))

        fig = go.Figure([trace0, trace1, trace2], layout=layout)

        iplot(fig, filename='pump-conditions.html')
```

```
In [8]: counts/counts.id.sum()
```

```
Out [8]:
```

	id
status_group	
functional	0.543081
functional needs repair	0.072677
non functional	0.384242

54.3% of pumps are functional, while 7.3% are functional but require repair and 38.4% are non functional.

1.2.2 Features

```
In [9]: train = train.set_index('id')
        train.head()
```

```
Out [9]:
```

	amount_tsh	date_recorded	funder	gps_height	installer	\
id						
69572	6000.0	2011-03-14	Roman	1390	Roman	
8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	
34310	25.0	2013-02-25	Lottery Club	686	World vision	
67743	0.0	2013-01-28	Unicef	263	UNICEF	
19728	0.0	2011-07-13	Action In A	0	Artisan	

	longitude	latitude	wpt_name	num_private	\
id					
69572	34.938093	-9.856322	none	0	

8776	34.698766	-2.147466		Zahanati	0
34310	37.460664	-3.821329		Kwa Mahundi	0
67743	38.486161	-11.155298	Zahanati	Ya Nanyumbu	0
19728	31.130847	-1.825359		Shuleni	0

	basin	...	payment_type	\
id		...		
69572	Lake Nyasa	...	annually	
8776	Lake Victoria	...	never pay	
34310	Pangani	...	per bucket	
67743	Ruvuma / Southern Coast	...	never pay	
19728	Lake Victoria	...	never pay	

	water_quality	quality_group	quantity	quantity_group	\
id					
69572	soft	good	enough	enough	
8776	soft	good	insufficient	insufficient	
34310	soft	good	enough	enough	
67743	soft	good	dry	dry	
19728	soft	good	seasonal	seasonal	

	source	source_type	source_class	\
id				
69572	spring	spring	groundwater	
8776	rainwater harvesting	rainwater harvesting	surface	
34310	dam	dam	surface	
67743	machine dbh	borehole	groundwater	
19728	rainwater harvesting	rainwater harvesting	surface	

	waterpoint_type	waterpoint_type_group
id		
69572	communal standpipe	communal standpipe
8776	communal standpipe	communal standpipe
34310	communal standpipe multiple	communal standpipe
67743	communal standpipe multiple	communal standpipe
19728	communal standpipe	communal standpipe

[5 rows x 39 columns]

In [10]: train.shape

Out[10]: (59400, 39)

The shape of the full feature data is (59400,39). Having more than 2 orders of magnitude worth of samples compared to the number of features will help avoid the curse of dimensionality. The concern would be with the label with the smallest sample size. The data for the condition "functional needs repair" is (4317, 39). In this case, the number of samples is still 2 orders of magnitude larger than the number of features.

The preview above shows that the training data contains 39 features with mixed datatypes. The descriptions of each feature are described below, from the data source.

amount_tsh - Total static head (amount water available to waterpoint) date_recorded - The date the row was entered funder - Who funded the well gps_height - Altitude of the well installer - Organization that installed the well longitude - GPS coordinate latitude - GPS coordinate wpt_name - Name of the waterpoint if there is one num_private - basin - Geographic water basin subvillage - Geographic location region - Geographic location region_code - Geographic location (coded) district_code - Geographic location (coded) lga - Geographic location ward - Geographic location population - Population around the well public_meeting - True/False recorded_by - Group entering this row of data scheme_management - Who operates the waterpoint scheme_name - Who operates the waterpoint permit - If the waterpoint is permitted construction_year - Year the waterpoint was constructed extraction_type - The kind of extraction the waterpoint uses extraction_type_group - The kind of extraction the waterpoint uses extraction_type_class - The kind of extraction the waterpoint uses management - How the waterpoint is managed management_group - How the waterpoint is managed payment - What the water costs payment_type - What the water costs water_quality - The quality of the water quality_group - The quality of the water quantity - The quantity of water quantity_group - The quantity of water source - The source of the water source_type - The source of the water source_class - The source of the water waterpoint_type - The kind of waterpoint waterpoint_type_group - The kind of waterpoint

Some of the feature descriptions look very similar to each other. I will be checking if there are duplicate features, or if some of the features are derived from others, this will help reduce the number of features to just the essential ones.

```
In [89]: def data_check(data, col):
          print('There are %s ' % data[col].isnull().sum() + ' null values.')
          print('There are %s ' % data[col].nunique() + ' unique values.')
```

```
In [90]: def plot_hist(data, col, ylog=False, xlog=False):

          if ylog:
              ymode='log'
          else:
              ymode=None
          if xlog:
              xmode='log'
          else:
              xmode=None

          trace = go.Histogram(x=data[col], name='col')

          layout = go.Layout(title=f'{col} Distribution',
                              yaxis=dict(title='Count', type=ymode),
                              xaxis=dict(type=xmode))

          fig = go.Figure([trace], layout=layout)

          iplot(fig, filename=f'{col}-dist.html')
```

amount_tsh amount_tsh is the "total static head" of the pump, described as the total amount of water available to the waterpoint. Total static head is the distance from the surface of the water source to the surface where the water is collected from the pump. It is related to the pressure of the water source and therefore the amount of water present. The distribution of amount_tsh is shown below on a semi-log plot. Most of the pumps are near 0 total static head.

```
In [12]: plot_hist(train, col='amount_tsh', ylog=True)
```

```
In [13]: train.amount_tsh.value_counts().head()
```

```
Out[13]: 0.0      41639
         500.0    3102
         50.0     2472
        1000.0    1488
         20.0     1463
         Name: amount_tsh, dtype: int64
```

The top 5 value counts for amount_tsh are summarized above. 41,639 of the 59,400 pumps have a total static head of 0. This implies the water source is at the surface.

```
In [14]: train.amount_tsh.isnull().any()
```

```
Out[14]: False
```

There are no null values in this feature column!

date_recorded

```
In [15]: train.date_recorded.isnull().any()
```

```
Out[15]: False
```

The date_recorded column also has no missing values.

```
In [16]: plot_hist(train, col='date_recorded')
```

The first date we have a record for is 2002-10-14. The data is very sparse prior to 2011. If we zoom in to Jan 2011 and beyond, there appear to be 4 periods of heavy data collection. After April 2013, there also appears to be some more regular data collection each month, on about 200-300 pumps at a time.

funder

```
In [91]: data_check(train, 'funder')
```

There are 3635 null values.

There are 1897 unique values.

3,635 of the 59,400 samples are missing their funding source. I can't imagine how the funding source would be related to a pump failing or not so I think it's safe to drop this column.

gps_height

```
In [19]: plot_hist(train, col='gps_height')
```

```
In [20]: train.gps_height.isnull().any()
```

```
Out[20]: False
```

The `gps_height` is the altitude of the water pump. Luckily, there are no missing data in this column.

installer

```
In [93]: data_check(train, 'installer')
```

```
There are 3655 null values.
```

```
There are 2145 unique values.
```

There are 2,145 unique installers. Unfortunately, there are also 3,655 missing values for `installer`. This feature can affect water pump condition if for example, the quality of the installation varies between installers.

latitude, longitude

```
In [94]: data_check(train, 'latitude')
```

```
There are 0 null values.
```

```
There are 57517 unique values.
```

```
In [95]: data_check(train, 'longitude')
```

```
There are 0 null values.
```

```
There are 57516 unique values.
```

No missing values for latitude or longitude!

```
In [23]: px.set_mapbox_access_token(credentials['mapbox_token'])
```

```
fig = px.scatter_mapbox(train, lat="latitude", lon="longitude",
                        color='gps_height',
                        size_max=15, zoom=4.5)
fig.show()
```

The map above shows the locations of the pumps with a color gradient showing the altitude of each pump (`gps_height`). The pumps are mostly in concentrated areas but there are some small groups of dispersed pump sites.

wpt_name

```
In [96]: data_check(train, 'wpt_name')
```

There are 0 null values.

There are 37400 unique values.

The wpt_name is simply the name of the water point. Interestingly, every waterpoint has a name, there are no null values. There are 37,400 unique names while we have 59,400 water points, so waterpoints can have duplicate names. Below are 20 sample names.

```
In [48]: train.wpt_name.unique()[:20]
```

```
Out[48]: array(['none', 'Zahanati', 'Kwa Mahundi', 'Zahanati Ya Nanyumbu',  
                'Shuleni', 'Tajiri', 'Kwa Ngomho', 'Tushirikiane',  
                'Kwa Ramadhan Musa', 'Kwapeto', 'Mzee Hokororo',  
                'Kwa Alid Nchimbi', 'Pamba', 'Kwa John Izack Mmari', 'Mwabasabi',  
                'Kwa Juvenal Ching'Ombe', 'Kwa John Mtenzi', 'Kwa Rose Chaula',  
                'Ngomee', 'Muungano'], dtype=object)
```

num_private

```
In [97]: data_check(train, 'num_private')
```

There are 0 null values.

There are 65 unique values.

```
In [54]: train.num_private.value_counts()
```

```
Out[54]: 0          58643  
         6           81  
         1           73  
         5           46  
         8           46  
        32           40  
        45           36  
        15           35  
        39           30  
        93           28  
         3           27  
         7           26  
         2           23  
        65           22  
        47           21  
       102           20  
         4           20  
        17           17
```


80	15
20	14
25	12
11	11
41	10
34	10
16	8
120	7
150	6
22	6
12	5
24	5
	...
14	3
61	3
27	2
26	2
160	1
30	1
698	1
60	1
1402	1
450	1
668	1
131	1
35	1
672	1
42	1
136	1
87	1
300	1
280	1
141	1
62	1
111	1
240	1
1776	1
755	1
180	1
213	1
23	1
55	1
94	1

Name: num_private, Length: 65, dtype: int64

Unfortunately, there is no description of num_private. There are 65 unique values for this feature with no missing data. The vast majority of samples however have a value of 0 for this feature, 58,643 out of 59,400 have a value of 0.

basin

```
In [98]: data_check(train, 'basin')
```

```
There are 0 null values.  
There are 9 unique values.
```

```
In [55]: plot_hist(train, col='basin')
```

basin refers to the water basin for the water source. The water basin with the most water pumps is Lake Victoria and the one with the least is Lake Rukwa. Below we can visualize the distribution of waterpumps according to which basin they draw water from. Lake Victoria is at the very North end of Tanzania, and we can see all of the neighboring waterpoints that access it.

```
In [63]: px.set_mapbox_access_token(credentials['mapbox_token'])
```

```
fig = px.scatter_mapbox(train, lat="latitude", lon="longitude",  
                        color='basin',  
                        size_max=15, zoom=4.5)  
  
fig.show()
```

subvillage

```
In [99]: data_check(train, 'subvillage')
```

```
There are 371 null values.  
There are 19287 unique values.
```

There are 19,287 unique subvillages with 371 missing values.

region

```
In [100]: data_check(train, 'region')
```

```
There are 0 null values.  
There are 21 unique values.
```

```
In [62]: plot_hist(train, col='region')
```

There are 21 unique regions with 0 missing values. We can see how there are multiple regions that draw from the same basins.

```
In [64]: px.set_mapbox_access_token(credentials['mapbox_token'])
```

```
fig = px.scatter_mapbox(train, lat="latitude", lon="longitude",  
                        color='region',  
                        size_max=15, zoom=4.5)  
  
fig.show()
```

region_code

```
In [101]: data_check(train, 'region_code')
```

There are 0 null values.

There are 27 unique values.

There are actually 27 unique region_codes which is unexpected. Let's see how they are mapped to region names.

```
In [70]: train[['date_recorded', 'region', 'region_code']].groupby(['region', 'region_code']).
```

```
Out[70]:
```

		date_recorded
--	--	---------------

region	region_code	date_recorded
Arusha	2	3024
	24	326
Dar es Salaam	7	805
Dodoma	1	2201
Iringa	11	5294
Kagera	18	3316
Kigoma	16	2816
Kilimanjaro	3	4379
Lindi	8	300
	18	8
	80	1238
Manyara	21	1583
Mara	20	1969
Mbeya	12	4639
Morogoro	5	4006
Mtwara	9	390
	90	917
	99	423
Mwanza	17	55
	19	3047
Pwani	6	1609
	40	1
	60	1025
Rukwa	15	1808
Ruvuma	10	2640
Shinyanga	11	6
	14	20
	17	4956
Singida	13	2093
Tabora	14	1959
Tanga	4	2513
	5	34

So a given region can have multiple region codes. region_codes are not completely unique across regions. Notice that region 5 appears both for the Tanga region and Morogoro.

district code

```
In [102]: data_check(train, 'district_code')
```

There are 0 null values.

There are 20 unique values.

There are 20 unique district codes with no missing values.

```
In [73]: train[['date_recorded', 'region_code', 'district_code']].groupby(['region_code', 'dis
```

```
Out[73]:
```

	region_code	district_code	date_recorded
1	0	0	23
		1	888
		3	361
		4	347
		5	358
		6	224
2	1	1	189
		2	1206
		3	109
		5	201
		6	310
		7	1009
3	1	1	595
		2	519
		3	877
		4	1225
		5	620
		6	109
		7	434
4	1	1	698
		2	408
		3	323
		4	110
		5	293
		6	266
		7	127
		8	288
5	1	1	1128
		2	521
		3	997
...			...
19	5	5	332
		6	488
		7	347
		8	132

20	1	171
	2	716
	3	396
	4	438
	6	248
21	1	550
	2	274
	3	297
	4	276
	5	186
24	30	326
40	43	1
60	33	115
	43	350
	53	454
	60	63
	63	37
	67	6
	80	13
23		293
43		154
53		291
62		109
90	33	759
	63	158
99	1	423

[130 rows x 1 columns]

The groupby table above shows that for each region_code, there can be multiple districts. The district codes are not unique between regions.

lga

```
In [103]: data_check(train, 'lga')
```

There are 0 null values.

There are 125 unique values.

lga has 125 unique values with no missing data.

```
In [75]: train[['date_recorded', 'region_code', 'district_code', 'lga']].groupby(['region_code
```

```
Out[75]:
```

	region_code	lga	district_code	date_recorded
1		Bahi	6	224
		Chamwino	4	347
		Dodoma Urban	5	358

	Kondoa	1	523
	Kongwa	3	361
	Mpwapwa	0	23
		1	365
2	Arusha Rural	2	1206
		3	46
	Arusha Urban	3	63
	Longido	6	310
	Meru	7	1009
	Monduli	1	189
	Ngorongoro	5	201
3	Hai	5	617
		6	8
	Moshi Rural	1	7
		4	1219
		5	3
		6	22
	Moshi Urban	6	79
	Mwanga	2	519
	Rombo	1	588
		4	6
	Same	3	877
	Siha	7	434
4	Handeni	6	254
	Kilindi	7	127
	Korogwe	1	4
		2	408
...			...
20	Musoma Rural	3	396
	Rorya	6	210
	Serengeti	2	716
	Tarime	1	171
		6	38
21	Babati	1	511
	Hanang	2	274
	Kiteto	4	7
		5	186
	Mbulu	3	297
	Simanjiro	1	39
		4	269
24	Karatu	30	326
40	Mkuranga	43	1
60	Kisarawe	33	115
	Mafia	60	63
		63	37
		67	6
	Mkuranga	43	350
	Rufiji	53	454

80	Kilwa	13	391
	Lindi Rural	23	293
		62	88
	Lindi Urban	62	21
	Liwale	43	154
	Ruangwa	53	291
90	Masasi	33	528
	Nanyumbu	63	158
	Newala	33	231
99	Mtwara Rural	1	423

[176 rows x 1 columns]

```
In [42]: plot_hist(train, col='lga')
```

ward

```
In [104]: data_check(train, 'ward')
```

There are 0 null values.

There are 2092 unique values.

ward has 2092 unique values with no missing data.

population

```
In [105]: data_check(train, 'population')
```

There are 0 null values.

There are 1049 unique values.

```
In [83]: plot_hist(train, col='population', ylog=True)
```

There are 1049 unique values of population with 0 missing data. This number is different than the unique number of water pumps, suggesting that multiple pumps can serve the same population, or there are multiple populations with the same number of people.

public meeting

```
In [109]: data_check(train, 'public_meeting')
```

There are 3334 null values.

There are 2 unique values.

```
In [107]: train.public_meeting.value_counts()
```

```
Out[107]: True      51011
          False     5055
          Name: public_meeting, dtype: int64
```

public_meeting is a binary feature. 3,334 values are missing for this feature.
No description is provided for this feature.

recorded_by

```
In [110]: data_check(train, 'recorded_by')
```

```
There are 0 null values.
```

```
There are 1 unique values.
```

```
In [111]: train.recorded_by.unique()
```

```
Out[111]: array(['GeoData Consultants Ltd'], dtype=object)
```

All of the data was recorded by GeoData Consultants. This value is the same for the entire dataset and thus can be dropped.

scheme_management

```
In [112]: data_check(train, 'scheme_management')
```

```
There are 3877 null values.
```

```
There are 12 unique values.
```

```
In [113]: plot_hist(train, col='scheme_management')
```

scheme_name

```
In [115]: data_check(train, 'scheme_name')
```

```
There are 28166 null values.
```

```
There are 2696 unique values.
```

permit

```
In [116]: data_check(train, 'permit')
```

```
There are 3056 null values.
```

```
There are 2 unique values.
```

```
In [117]: train.permit.value_counts()
```

```
Out[117]: True      38852  
          False    17492  
          Name: permit, dtype: int64
```


construction_year

```
In [118]: data_check(train, 'construction_year')
```

There are 0 null values.

There are 55 unique values.

```
In [121]: train.construction_year.value_counts()
```

```
Out[121]: 0          20709
          2010          2645
          2008          2613
          2009          2533
          2000          2091
          2007          1587
          2006          1471
          2003          1286
          2011          1256
          2004          1123
          2012          1084
          2002          1075
          1978          1037
          1995          1014
          2005          1011
          1999           979
          1998           966
          1990           954
          1985           945
          1980           811
          1996           811
          1984           779
          1982           744
          1994           738
          1972           708
          1974           676
          1997           644
          1992           640
          1993           608
          2001           540
          1988           521
          1983           488
          1975           437
          1986           434
          1976           414
          1970           411
          1991           324
          1989           316
          1987           302
```

1981	238
1977	202
1979	192
1973	184
2013	176
1971	145
1960	102
1967	88
1963	85
1968	77
1969	59
1964	40
1962	30
1961	21
1965	19
1966	17

Name: construction_year, dtype: int64

There are no null values for construction_year, however 20,709 of the waterpoints have a construction year of 0. This likely means that the year of construction year is not known for the water pump.

extraction_type

```
In [122]: data_check(train, 'extraction_type')
```

There are 0 null values.
There are 18 unique values.

```
In [123]: plot_hist(train, col='extraction_type')
```

There are 18 unique extraction types with no missing data. The most common is the gravity extraction type with 26,780 water pumps of this type.

extraction_type_group

```
In [124]: data_check(train, 'extraction_type_group')
```

There are 0 null values.
There are 13 unique values.

```
In [125]: plot_hist(train, col='extraction_type_group')
```

extraction_type_class

```
In [126]: data_check(train, 'extraction_type_class')
```

There are 0 null values.
There are 7 unique values.

```
In [127]: plot_hist(train, col='extraction_type_class')
```

```
In [131]: train[['extraction_type_class', 'extraction_type_group', 'extraction_type', 'date_re
```

```
Out[131]:
```

extraction_type_class	extraction_type_group	extraction_type	date_recorded
gravity	gravity	gravity	26780
handpump	afridev	afridev	1770
	india mark ii	india mark ii	2400
	india mark iii	india mark iii	98
	nira/tanira	nira/tanira	8154
	other handpump	other - mkulima/shinyanga	2
		other - play pump	85
		other - swn 81	229
		walimi	48
	swn 80	swn 80	3670
motorpump	mono	mono	2865
	other motorpump	chemo	90
		climax	32
other	other	other	6430
rope pump	rope pump	other - rope pump	451
submersible	submersible	ksb	1415
		submersible	4764
wind-powered	wind-powered	windmill	117

The groupby above shows the breakdown from extraction class to group to type. Since class and group add no new information that's not already covered by extraction_type, we can drop these two extra features.

management

```
In [129]: data_check(train, 'management')
```

There are 0 null values.
There are 12 unique values.

```
In [128]: plot_hist(train, col='management')
```

management_group

```
In [132]: data_check(train, 'management_group')
```

There are 0 null values.
There are 5 unique values.

```
In [133]: plot_hist(train, col='management_group')
```

```
In [170]: train[['management_group', 'management', 'date_recorded']].groupby(['management_group', 'management']).count()
```

```
Out[170]:
```

management_group	management	date_recorded
commercial	company	685
	private operator	1971
	trust	78
	water authority	904
other	other	844
	other - school	99
	parastatal	1768
parastatal	parastatal	1768
unknown	unknown	561
user-group	vwc	40507
	water board	2933
	wua	2535
	wug	6515

payment

```
In [137]: data_check(train, 'payment')
```

There are 0 null values.

There are 7 unique values.

```
In [138]: plot_hist(train, col='payment')
```

payment_type

```
In [139]: data_check(train, 'payment_type')
```

There are 0 null values.

There are 7 unique values.

```
In [140]: plot_hist(train, col='payment_type')
```

```
In [142]: train[['payment_type', 'payment', 'date_recorded']].groupby(['payment_type', 'payment']).count()
```

```
Out[142]:
```

payment_type	payment	date_recorded
annually	pay annually	3642
monthly	pay monthly	8300
never pay	never pay	25348
on failure	pay when scheme fails	3914
other	other	1054
per bucket	pay per bucket	8985
unknown	unknown	8157

payment_type and payment are duplicate columns. We can drop one of them.

water_quality

```
In [144]: data_check(train, 'water_quality')
```

There are 0 null values.

There are 8 unique values.

```
In [145]: plot_hist(train, col='water_quality')
```

quality_group

```
In [146]: data_check(train, 'quality_group')
```

There are 0 null values.

There are 6 unique values.

```
In [147]: plot_hist(train, col='quality_group')
```

```
In [149]: train[['quality_group', 'water_quality', 'date_recorded']].groupby(['quality_group',
```

```
Out[149]:
```

		date_recorded
quality_group	water_quality	
colored	coloured	490
fluoride	fluoride	200
	fluoride abandoned	17
good	soft	50818
milky	milky	804
salty	salty	4856
	salty abandoned	339
unknown	unknown	1876

Again, quality_group gives us no new information so it can be dropped.

quantity

```
In [151]: data_check(train, 'quantity')
```

There are 0 null values.

There are 5 unique values.

```
In [152]: plot_hist(train, col='quantity')
```

quantity_group

```
In [153]: data_check(train, 'quantity_group')
```

There are 0 null values.

There are 5 unique values.

```
In [155]: train[['quantity_group', 'quantity', 'date_recorded']].groupby(['quantity_group', 'q
```

```
Out[155]:
```

		date_recorded
quantity_group	quantity	
dry	dry	6246
enough	enough	33186
insufficient	insufficient	15129
seasonal	seasonal	4050
unknown	unknown	789

```
In [173]: (train.quantity_group==train.quantity).all()
```

```
Out[173]: True
```

quantity_group and quantity are also duplicate features, we can drop quantity_group.

source

```
In [156]: data_check(train, 'source')
```

There are 0 null values.

There are 10 unique values.

```
In [157]: plot_hist(train, col='source')
```

source_type

```
In [158]: data_check(train, 'source_type')
```

There are 0 null values.

There are 7 unique values.

```
In [159]: plot_hist(train, col='source_type')
```

source_class

```
In [161]: data_check(train, 'source_class')
```

There are 0 null values.

There are 3 unique values.

```
In [162]: plot_hist(train, col='source_class')
```

```
In [163]: train[['source_class', 'source_type', 'source', 'date_recorded']].groupby(['source_c
```

```
Out[163]:
```

	source_class	source_type	source	date_recorded
	groundwater	borehole	hand dtw	874
			machine dbh	11075
		shallow well	shallow well	16824
		spring	spring	17021
	surface	dam	dam	656
		rainwater harvesting	rainwater harvesting	2295
		river/lake	lake	765
			river	9612
	unknown	other	other	212
			unknown	66

waterpoint_type

```
In [165]: data_check(train, 'waterpoint_type')
```

There are 0 null values.

There are 7 unique values.

```
In [166]: plot_hist(train, col='waterpoint_type')
```

waterpoint_type_group

```
In [167]: data_check(train, 'waterpoint_type_group')
```

There are 0 null values.

There are 6 unique values.

```
In [168]: plot_hist(train, col='waterpoint_type_group')
```

```
In [169]: train[['waterpoint_type_group', 'waterpoint_type', 'date_recorded']].groupby(['waterp
```

```

Out[169]:

```

waterpoint_type_group	waterpoint_type	date_recorded
cattle trough	cattle trough	116
communal standpipe	communal standpipe	28522
	communal standpipe multiple	6103
dam	dam	7
hand pump	hand pump	17488
improved spring	improved spring	784
other	other	6380

At this point we saw several features that are related to each other and are broken down into increasing granularity. We saw this for `waterpoint_type`, `source`, `water_quality`, `extraction_type`, and `management`. It's unclear to me whether the machine learning models would perform best if I was to keep only the most granular features and drop the higher categories. My plan is to create two datasets, one that keeps all of the levels of these features and one that keeps only the most granular and compare their performance.