

Capstone 2 - Predicting Water Pump Condition in Tanzania Data Munging

September 8, 2019

1 Capstone 2 - Predicting Water Pump Condition in Tanzania Data Munging

Kenneth Liao

1.1 Background

The UN publishes and reviews a list of least developed countries (LDC) every 3 years. LDCs are “low-income countries confronting severe structural impediments to sustainable development. They are highly vulnerable to economic and environmental shocks and have low levels of human assets.”¹. Tanzania has been classified as an LDC since the UN published the first list of LDCs in 1971². A common challenge of LDCs is a lack of infrastructure to support the development of the nation, including access to education and healthcare, waste management, and access to potable water.

According to UNICEF, as of 2017, more than 24 million Tanzanians lacked access to basic drinking water³. This corresponds to only 56.7% of the country’s population having access to basic drinking water. Outside of developed urban areas, much of the potable water is accessed via water pumps.

Taarifa is an open-source platform for crowd-sourced reporting and triaging of infrastructure related issues. Together with the Tanzanian Ministry of Water, data has been collected for thousands of water pumps throughout Tanzania. The goal of this project is to be able to predict the condition of these water pumps to improve maintenance, reduce pump downtime, and ensure basic water access for millions of Tanzanians.

References

1. <https://www.un.org/development/desa/dpad/least-developed-country-category.html>
2. https://www.un.org/development/desa/dpad/wp-content/uploads/sites/45/publication/ldc_list.pdf
3. <https://washwatch.org/en/countries/tanzania/summary/statistics/>

1.1.1 Problem Description

Predict the operating condition of water pumps in Tanzania given various metadata on each water pump.

1.1.2 Strategy

The strategy will be to implement an XGBoost model as well as a neural network model for predictions and compare their performance.

1.1.3 Data

The dataset is provided by Taarifa, together with the Tanzanian Ministry of Water and is hosted by DrivenData.org:

<https://www.drivendata.org/competitions/7/pump-it-up-data-mining-the-water-table/page/23/>

1.2 Data Munging

```
In [1]: import pandas as pd
import plotly.graph_objs as go
from plotly.offline import iplot, plot, init_notebook_mode
import plotly.express as px
from config import credentials

init_notebook_mode(connected=True)
```

```
In [2]: # load the data
train = pd.read_csv('../data/train.csv')
train_labels = pd.read_csv('../data/train-labels.csv')
```

I'll start by removing the unwanted feature columns we identified in the EDA part of the analysis. This includes duplicate, irrelevant, and single value columns.

```
In [3]: duplicated = ['recorded_by', 'payment_type', 'quantity_group']

train_clean = train.drop(duplicated, axis=1)
train_clean.columns
```

```
Out[3]: Index(['id', 'amount_tsh', 'date_recorded', 'funder', 'gps_height',
              'installer', 'longitude', 'latitude', 'wpt_name', 'num_private',
              'basin', 'subvillage', 'region', 'region_code', 'district_code', 'lga',
              'ward', 'population', 'public_meeting', 'scheme_management',
              'scheme_name', 'permit', 'construction_year', 'extraction_type',
              'extraction_type_group', 'extraction_type_class', 'management',
              'management_group', 'payment', 'water_quality', 'quality_group',
              'quantity', 'source', 'source_type', 'source_class', 'waterpoint_type',
              'waterpoint_type_group'],
              dtype='object')
```

```
In [4]: train_clean.set_index(['id', 'date_recorded'], inplace=True)
```

```
In [5]: train_clean.head()
```

Out [5]:

		amount_tsh	funder	gps_height	installer	\
id	date_recorded					
69572	2011-03-14	6000.0	Roman	1390	Roman	
8776	2013-03-06	0.0	Grumeti	1399	GRUMETI	
34310	2013-02-25	25.0	Lottery Club	686	World vision	
67743	2013-01-28	0.0	Unicef	263	UNICEF	
19728	2011-07-13	0.0	Action In A	0	Artisan	

		longitude	latitude	wpt_name	num_private	\
id	date_recorded					
69572	2011-03-14	34.938093	-9.856322	none	0	
8776	2013-03-06	34.698766	-2.147466	Zahanati	0	
34310	2013-02-25	37.460664	-3.821329	Kwa Mahundi	0	
67743	2013-01-28	38.486161	-11.155298	Zahanati Ya Nanyumbu	0	
19728	2011-07-13	31.130847	-1.825359	Shuleni	0	

		basin	subvillage	\
id	date_recorded			
69572	2011-03-14	Lake Nyasa	Mnyusi B	
8776	2013-03-06	Lake Victoria	Nyamara	
34310	2013-02-25	Pangani	Majengo	
67743	2013-01-28	Ruvuma / Southern Coast	Mahakamani	
19728	2011-07-13	Lake Victoria	Kyanyamisa	

		...	management_group	payment	\
id	date_recorded	...			
69572	2011-03-14	...	user-group	pay annually	
8776	2013-03-06	...	user-group	never pay	
34310	2013-02-25	...	user-group	pay per bucket	
67743	2013-01-28	...	user-group	never pay	
19728	2011-07-13	...	other	never pay	

		water_quality	quality_group	quantity	\
id	date_recorded				
69572	2011-03-14	soft	good	enough	
8776	2013-03-06	soft	good	insufficient	
34310	2013-02-25	soft	good	enough	
67743	2013-01-28	soft	good	dry	
19728	2011-07-13	soft	good	seasonal	

		source	source_type	source_class	\
id	date_recorded				
69572	2011-03-14	spring	spring	groundwater	
8776	2013-03-06	rainwater harvesting	rainwater harvesting	surface	
34310	2013-02-25	dam	dam	surface	
67743	2013-01-28	machine dbh	borehole	groundwater	
19728	2011-07-13	rainwater harvesting	rainwater harvesting	surface	

id	date_recorded	waterpoint_type	waterpoint_type_group
69572	2011-03-14	communal standpipe	communal standpipe
8776	2013-03-06	communal standpipe	communal standpipe
34310	2013-02-25	communal standpipe multiple	communal standpipe
67743	2013-01-28	communal standpipe multiple	communal standpipe
19728	2011-07-13	communal standpipe	communal standpipe

[5 rows x 35 columns]

Next, I need to convert the categorical text features into dummy variables.

```
In [6]: # list of all categorical variables
cat_cols = []
for col in train_clean.columns:
    if train_clean[col].dtype == 'object':
        cat_cols.append(col)
cat_cols
```

```
Out[6]: ['funder',
         'installer',
         'wpt_name',
         'basin',
         'subvillage',
         'region',
         'lga',
         'ward',
         'public_meeting',
         'scheme_management',
         'scheme_name',
         'permit',
         'extraction_type',
         'extraction_type_group',
         'extraction_type_class',
         'management',
         'management_group',
         'payment',
         'water_quality',
         'quality_group',
         'quantity',
         'source',
         'source_type',
         'source_class',
         'waterpoint_type',
         'waterpoint_type_group']
```

```
In [7]: %%time
cat_dummies = pd.get_dummies(train_clean[cat_cols], dummy_na=True)
```

Wall time: 1min 20s

I use `pd.get_dummies` with the argument `dummy_na=True` so that null values are not ignored. They are instead encoded the same as all other values so each feature will have a null dummy variable, indicated whether the sample was null or not for that feature. The resulting categorical feature set now has 65,828 features.

```
In [8]: cat_dummies.head()
```

```
Out[8]:
```

		funder_0	funder_A/co	Germany	funder_Aar	\
id	date_recorded					
69572	2011-03-14	0		0		0
8776	2013-03-06	0		0		0
34310	2013-02-25	0		0		0
67743	2013-01-28	0		0		0
19728	2011-07-13	0		0		0

		funder_Abas	Ka	funder_Abasia	\
id	date_recorded				
69572	2011-03-14		0	0	
8776	2013-03-06		0	0	
34310	2013-02-25		0	0	
67743	2013-01-28		0	0	
19728	2011-07-13		0	0	

		funder_Abc-ihushi	Development	Cent	funder_Abd	\
id	date_recorded					
69572	2011-03-14			0		0
8776	2013-03-06			0		0
34310	2013-02-25			0		0
67743	2013-01-28			0		0
19728	2011-07-13			0		0

		funder_Abdala	funder_Abddwe	funder_Abdul	\
id	date_recorded				
69572	2011-03-14	0	0	0	
8776	2013-03-06	0	0	0	
34310	2013-02-25	0	0	0	
67743	2013-01-28	0	0	0	
19728	2011-07-13	0	0	0	

		...	\
id	date_recorded	...	
69572	2011-03-14	...	
8776	2013-03-06	...	
34310	2013-02-25	...	
67743	2013-01-28	...	

19728 2011-07-13

...

id	date_recorded	waterpoint_type_improved spring	waterpoint_type_other \
69572	2011-03-14	0	0
8776	2013-03-06	0	0
34310	2013-02-25	0	0
67743	2013-01-28	0	0
19728	2011-07-13	0	0

id	date_recorded	waterpoint_type_nan	waterpoint_type_group_cattle trough \
69572	2011-03-14	0	0
8776	2013-03-06	0	0
34310	2013-02-25	0	0
67743	2013-01-28	0	0
19728	2011-07-13	0	0

id	date_recorded	waterpoint_type_group_communal standpipe \
69572	2011-03-14	1
8776	2013-03-06	1
34310	2013-02-25	1
67743	2013-01-28	1
19728	2011-07-13	1

id	date_recorded	waterpoint_type_group_dam \
69572	2011-03-14	0
8776	2013-03-06	0
34310	2013-02-25	0
67743	2013-01-28	0
19728	2011-07-13	0

id	date_recorded	waterpoint_type_group_hand pump \
69572	2011-03-14	0
8776	2013-03-06	0
34310	2013-02-25	0
67743	2013-01-28	0
19728	2011-07-13	0

id	date_recorded	waterpoint_type_group_improved spring \
69572	2011-03-14	0
8776	2013-03-06	0
34310	2013-02-25	0
67743	2013-01-28	0

19728 2011-07-13

0

		waterpoint_type_group_other	waterpoint_type_group_nan
id	date_recorded		
69572	2011-03-14	0	0
8776	2013-03-06	0	0
34310	2013-02-25	0	0
67743	2013-01-28	0	0
19728	2011-07-13	0	0

[5 rows x 65828 columns]

```
In [9]: # list of all numerical variables
num_cols = []
for col in train_clean.columns:
    if train_clean[col].dtype != 'object':
        num_cols.append(col)
num_cols
```

```
Out[9]: ['amount_tsh',
'gps_height',
'longitude',
'latitude',
'num_private',
'region_code',
'district_code',
'population',
'construction_year']
```

```
In [10]: numerical = train_clean[num_cols]
numerical.head()
```

```
Out[10]:
```

		amount_tsh	gps_height	longitude	latitude	\
id	date_recorded					
69572	2011-03-14	6000.0	1390	34.938093	-9.856322	
8776	2013-03-06	0.0	1399	34.698766	-2.147466	
34310	2013-02-25	25.0	686	37.460664	-3.821329	
67743	2013-01-28	0.0	263	38.486161	-11.155298	
19728	2011-07-13	0.0	0	31.130847	-1.825359	

		num_private	region_code	district_code	population	\
id	date_recorded					
69572	2011-03-14	0	11	5	109	
8776	2013-03-06	0	20	2	280	
34310	2013-02-25	0	21	4	250	
67743	2013-01-28	0	90	63	58	
19728	2011-07-13	0	18	1	0	

construction_year

id	date_recorded	
69572	2011-03-14	1999
8776	2013-03-06	2010
34310	2013-02-25	2009
67743	2013-01-28	1986
19728	2011-07-13	0

```
In [11]: numerical.info()
```

```
<class 'pandas.core.frame.DataFrame'>
MultiIndex: 59400 entries, (69572, 2011-03-14) to (26348, 2011-03-23)
Data columns (total 9 columns):
amount_tsh      59400 non-null float64
gps_height      59400 non-null int64
longitude       59400 non-null float64
latitude        59400 non-null float64
num_private     59400 non-null int64
region_code     59400 non-null int64
district_code   59400 non-null int64
population      59400 non-null int64
construction_year 59400 non-null int64
dtypes: float64(3), int64(6)
memory usage: 4.9+ MB
```

Luckily, none of the numerical columns have null values. We also don't need to normalize the numerical columns if using a tree-based model. However, for a neural network model, normalization will be necessary. I'll leave the data as-is for now and we can apply normalization when working with the neural network model specifically.

```
In [12]: # merge data back together.
```

```
In [13]: train_full = pd.concat([cat_dummies, numerical], axis=1)
train_full.head()
```

```
Out[13]:
```

		funder_0	funder_A/co	Germany	funder_Aar	\
id	date_recorded					
69572	2011-03-14	0		0		0
8776	2013-03-06	0		0		0
34310	2013-02-25	0		0		0
67743	2013-01-28	0		0		0
19728	2011-07-13	0		0		0

		funder_Abas	Ka	funder_Abasia	\
id	date_recorded				
69572	2011-03-14		0	0	
8776	2013-03-06		0	0	
34310	2013-02-25		0	0	
67743	2013-01-28		0	0	

19728	2011-07-13	0	0
-------	------------	---	---

		funder_Abc-ihushi Development Cent	funder_Abd \
id	date_recorded		
69572	2011-03-14	0	0
8776	2013-03-06	0	0
34310	2013-02-25	0	0
67743	2013-01-28	0	0
19728	2011-07-13	0	0

		funder_Abdala	funder_Abddwe	funder_Abdul \
id	date_recorded			
69572	2011-03-14	0	0	0
8776	2013-03-06	0	0	0
34310	2013-02-25	0	0	0
67743	2013-01-28	0	0	0
19728	2011-07-13	0	0	0

		...	waterpoint_type_group_nan	amount_tsh \
id	date_recorded	...		
69572	2011-03-14	...	0	6000.0
8776	2013-03-06	...	0	0.0
34310	2013-02-25	...	0	25.0
67743	2013-01-28	...	0	0.0
19728	2011-07-13	...	0	0.0

		gps_height	longitude	latitude	num_private \
id	date_recorded				
69572	2011-03-14	1390	34.938093	-9.856322	0
8776	2013-03-06	1399	34.698766	-2.147466	0
34310	2013-02-25	686	37.460664	-3.821329	0
67743	2013-01-28	263	38.486161	-11.155298	0
19728	2011-07-13	0	31.130847	-1.825359	0

		region_code	district_code	population	construction_year
id	date_recorded				
69572	2011-03-14	11	5	109	1999
8776	2013-03-06	20	2	280	2010
34310	2013-02-25	21	4	250	2009
67743	2013-01-28	90	63	58	1986
19728	2011-07-13	18	1	0	0

[5 rows x 65837 columns]

In [15]: train_full.info()

<class 'pandas.core.frame.DataFrame'>

MultiIndex: 59400 entries, (69572, 2011-03-14) to (26348, 2011-03-23)

```
Columns: 65837 entries, funder_0 to construction_year  
dtypes: float64(3), int64(6), uint8(65828)  
memory usage: 3.6+ GB
```

```
In [19]: train_full.to_pickle('../data/train_full.pkl')
```

The full dataset is now ready to train on. There may be issues with the dimension of this dataset after converting to dummy variables. The shape of the dataset is now 59400 X 69572. If the model shows poor performance, it may benefit by using another model to reduce the number of features to those which are most important. This can be done with a number of techniques including PCA, step-wise feature selection, and genetic algorithms for feature selection.